

Assignment Part A

Part I:

We have run our code for three different values of n, i.e N = 1024, 2048, 4096.

Unoptimised Code:

In the unoptimized code, we got to see with perf how the application is running and utilizing our resources.

N=1024,

Input matrix of size 1024

Reference execution time: 2286.64 ms

Performance counter stats for './cmmref data/input_1024.in':

4,595.58 msec task-clock	#	0.998 CPUs utilized	
276 context-switches	#	0.060 K/sec	
2 cpu-migrations	#	0.000 K/sec	
2,699 page-faults	#	0.587 K/sec	
18,19,48,12,353 cycles	#	3.959 GHz	(30.66%)
27,05,98,09,326 instructions	#	1.49 insn per cycle	(38.39%)
2,41,19,25,688 branches	#	524.836 M/sec	(38.57%)
32,64,541 branch-misses	#	0.14% of all branches	(38.66%)
14,34,12,75,416 L1-dcache-loads	#	3120.670 M/sec	(38.75%)
1,06,86,79,524 L1-dcache-load-misses	#	7.45% of all L1-dcache accesses	(38.76%)
1,06,47,11,857 LLC-loads	#	231.682 M/sec	(30.85%)
1,37,09,083 LLC-load-misses	#	1.29% of all LL-cache accesses	(30.76%)
<not supported> L1-icache-loads			
12,22,208 L1-icache-load-misses			(30.67%)
14,36,29,23,729 dTLB-loads	#	3125.380 M/sec	(30.58%)
68,692 dTLB-load-misses	#	0.00% of all dTLB cache accesses	(30.58%)
11,508 iTLB-loads	#	0.003 M/sec	(30.58%)
7,622 iTLB-load-misses	#	66.23% of all iTLB cache accesses	(30.58%)
<not supported> L1-dcache-prefetches			
<not supported> L1-dcache-prefetch-misses			

4.605288573 seconds time elapsed

4.588244000 seconds user

0.008035000 seconds sys

N=2048,

Input matrix of size 2048

Reference execution time: 63910.7 ms

Performance counter stats for './cmmref data/input_2048.in':

1,30,442.67 msec task-clock	#	1.000 CPUs utilized
-----------------------------	---	---------------------

612	context-switches	#	0.005 K/sec	
55	cpu-migrations	#	0.000 K/sec	
10,381	page-faults	#	0.080 K/sec	
5,18,96,67,80,573	cycles	#	3.979 GHz	(30.76%)
2,10,94,09,75,436	instructions	#	0.41 insn per cycle	(38.46%)
18,15,95,50,754	branches	#	139.215 M/sec	(38.46%)
1,45,35,344	branch-misses	#	0.08% of all branches	(38.46%)
1,13,24,77,29,533	L1-dcache-loads	#	868.180 M/sec	(38.47%)
13,71,94,19,159	L1-dcache-load-misses	#	12.11% of all L1-dcache accesses	(38.47%)
8,72,86,22,727	LLC-loads	#	66.915 M/sec	(30.77%)
6,56,71,48,703	LLC-load-misses	#	75.24% of all LL-cache accesses	(30.77%)
<not supported>	L1-icache-loads			
1,48,46,294	L1-icache-load-misses			(30.77%)
1,13,42,74,45,364	dTLB-loads	#	869.558 M/sec	(30.77%)
8,59,73,28,241	dTLB-load-misses	#	7.58% of all dTLB cache accesses	(30.77%)
1,26,983	iTLB-loads	#	0.973 K/sec	(30.76%)
3,79,312	iTLB-load-misses	#	298.71% of all iTLB cache accesses	(30.76%)
<not supported>	L1-dcache-prefetches			
<not supported>	L1-dcache-prefetch-misses			

130.449578975 seconds time elapsed

130.415606000 seconds user

0.027999000 seconds sys

N=4096,

Input matrix of size 4096

Reference execution time: 574311 ms

Performance counter stats for './cmmref data/input_4096.in':

11,44,275.46 msec	task-clock	#	1.000 CPUs utilized	
5,501	context-switches	#	0.005 K/sec	
256	cpu-migrations	#	0.000 K/sec	
41,099	page-faults	#	0.036 K/sec	
45,51,06,61,98,212	cycles	#	3.977 GHz	(30.77%)
16,69,80,17,01,832	instructions	#	0.37 insn per cycle	(38.46%)
1,41,65,87,77,199	branches	#	123.798 M/sec	(38.46%)
6,74,06,734	branch-misses	#	0.05% of all branches	(38.46%)
9,00,31,21,08,364	L1-dcache-loads	#	786.797 M/sec	(38.46%)
1,46,55,09,92,383	L1-dcache-load-misses	#	16.28% of all L1-dcache accesses	(38.46%)
70,79,74,25,908	LLC-loads	#	61.871 M/sec	(30.77%)
69,61,31,98,368	LLC-load-misses	#	98.33% of all LL-cache accesses	(30.77%)
<not supported>	L1-icache-loads			
12,24,46,694	L1-icache-load-misses			(30.77%)
8,99,96,96,86,505	dTLB-loads	#	786.497 M/sec	(30.77%)
68,74,87,23,435	dTLB-load-misses	#	7.64% of all dTLB cache accesses	(30.77%)
15,63,069	iTLB-loads	#	0.001 M/sec	(30.77%)
28,13,848	iTLB-load-misses	#	180.02% of all iTLB cache accesses	(30.77%)
<not supported>	L1-dcache-prefetches			
<not supported>	L1-dcache-prefetch-misses			

1144.320582699 seconds time elapsed

1144.189403000 seconds user

0.087996000 seconds sys

Basically we are observing that the L1- Data cache misses are high (7.45%, 12.11%, 16.28%). It is basically due to the access of different blocks of elements in the matrix for the unoptimised code. The unoptimised code is accessing the matrix element that are not currently on the L1 Data cache. Also we can clearly observe that the application takes most of the time in its implementation of checkered matrix multiplication.

Due to the Amdahl's Law, we come to know that if we apply optimisation on the matrix multiplication part the speedup is guaranteed.

Now we apply our optimised code and run the code.

Optimised Code:

Single Thread:

In the single thread, we basically have to improve the algorithm of the unoptimised code. So, in my optimisation, I have implemented the code by calculating the element of the new matrix serially and used the contiguous matrix elements of A and B to calculate so that the D cache misses are less. As the code uses the 1st row of A and 1st columnset of B. Here we see that if we implement B in column major, we have huge advantage because the code will use contiguous access of B.

So in single thread, I have converted B into column major order and calculated the elements of C serially and thus optimising our code.

N=1024,

Input matrix of size 1024

Reference execution time: 0 ms

Single thread execution time: 1332.31 ms

Performance counter stats for './cmmsingle data/input_1024.in':

3,403.17 msec task-clock	#	1.000 CPUs utilized	
12 context-switches	#	0.004 K/sec	
1 cpu-migrations	#	0.000 K/sec	
3,212 page-faults	#	0.944 K/sec	
13,54,38,13,577 cycles	#	3.980 GHz	(30.67%)
31,58,11,89,631 instructions	#	2.33 insn per cycle	(38.43%)
3,20,17,15,356 branches	#	940.804 M/sec	(38.55%)
25,24,699 branch-misses	#	0.08% of all branches	(38.66%)
15,96,91,39,609 L1-dcache-loads	#	4692.431 M/sec	(38.78%)
60,54,54,718 L1-dcache-load-misses	#	3.79% of all L1-dcache accesses	(38.77%)
53,85,95,687 LLC-loads	#	158.263 M/sec	(30.89%)
6,51,034 LLC-load-misses	#	0.12% of all LL-cache accesses	(30.78%)
<not supported> L1-icache-loads			
8,87,947 L1-icache-load-misses			(30.66%)
15,94,19,60,772 dTLB-loads	#	4684.445 M/sec	(30.56%)
12,447 dTLB-load-misses	#	0.00% of all dTLB cache accesses	(30.56%)
7,320 iTLB-loads	#	0.002 M/sec	(30.56%)
4,574 iTLB-load-misses	#	62.49% of all iTLB cache accesses	(30.56%)

<not supported> L1-dcache-prefetches
<not supported> L1-dcache-prefetch-misses

3.403526695 seconds time elapsed

3.399540000 seconds user

0.003999000 seconds sys

N=2048,

Input matrix of size 2048

Reference execution time: 0 ms

Single thread execution time: 11610.9 ms

Performance counter stats for './cmmsingle data/input_2048.in':

78,660.82 msec task-clock	#	1.000 CPUs utilized	
1,243 context-switches	#	0.016 K/sec	
32 cpu-migrations	#	0.000 K/sec	
12,428 page-faults	#	0.158 K/sec	
3,11,15,14,42,360 cycles	#	3.956 GHz	(30.77%)
2,47,48,04,48,755 instructions	#	0.80 insn per cycle	(38.46%)
24,61,80,98,786 branches	#	312.965 M/sec	(38.46%)
1,12,69,958 branch-misses	#	0.05% of all branches	(38.46%)
1,26,08,72,69,103 L1-dcache-loads	#	1602.923 M/sec	(38.46%)
7,41,79,02,528 L1-dcache-load-misses	#	5.88% of all L1-dcache accesses	(38.46%)
4,37,48,65,564 LLC-loads	#	55.617 M/sec	(30.77%)
3,30,11,39,766 LLC-load-misses	#	75.46% of all LL-cache accesses	(30.77%)
<not supported> L1-icache-loads			
1,28,83,895 L1-icache-load-misses			(30.77%)
1,26,27,61,08,045 dTLB-loads	#	1605.324 M/sec	(30.77%)
4,30,51,07,034 dTLB-load-misses	#	3.41% of all dTLB cache accesses	(30.77%)
1,64,268 iTLB-loads	#	0.002 M/sec	(30.76%)
2,25,261 iTLB-load-misses	#	137.13% of all iTLB cache accesses	(30.76%)
<not supported> L1-dcache-prefetches			
<not supported> L1-dcache-prefetch-misses			

78.692315373 seconds time elapsed

78.617746000 seconds user

0.043989000 seconds sys

N=4096,

Input matrix of size 4096

Reference execution time: 0 ms

Single thread execution time: 93343.7 ms

Performance counter stats for './cmmsingle data/input_4096.in':

95,034.11 msec task-clock	#	1.000 CPUs utilized	
503 context-switches	#	0.005 K/sec	
94 cpu-migrations	#	0.001 K/sec	
41,100 page-faults	#	0.432 K/sec	
3,64,44,44,38,543 cycles	#	3.835 GHz	(30.76%)

11,34,75,89,79,815	instructions	# 3.11 insn per cycle	(38.46%)
1,24,08,31,34,633	branches	# 1305.669 M/sec	(38.46%)
2,51,16,153	branch-misses	# 0.02% of all branches	(38.46%)
5,55,97,63,41,051	L1-dcache-loads	# 5850.282 M/sec	(38.47%)
6,63,75,64,728	L1-dcache-load-misses	# 1.19% of all L1-dcache accesses	(38.47%)
8,93,24,585	LLC-loads	# 0.940 M/sec	(30.77%)
7,48,11,698	LLC-load-misses	# 83.75% of all LL-cache accesses	(30.77%)
<not supported>	L1-icache-loads		
2,06,78,699	L1-icache-load-misses		(30.77%)
5,55,88,71,99,791	dTLB-loads	# 5849.344 M/sec	(30.77%)
7,62,14,693	dTLB-load-misses	# 0.01% of all dTLB cache accesses	(30.77%)
4,90,406	iTLB-loads	# 0.005 M/sec	(30.76%)
3,45,595	iTLB-load-misses	# 70.47% of all iTLB cache accesses	(30.76%)
<not supported>	L1-dcache-prefetches		
<not supported>	L1-dcache-prefetch-misses		

95.038998524 seconds time elapsed

94.963109000 seconds user

0.071999000 seconds sys

Multi-thread:

In the multi-thread implementation, I have basically used the implementation of single thread and divided the rows into some threads and calculated the matrix multiplication parallelly.

In the testing, we can see that for no of threads increasing from 1 to 2, the speedup is likely to be nearly 2x times. And if we make it 4, the speedup is almost 4x times for ideal but here close to 3. So in this way, we can show the scalability of our implementation.

N=1024,

Input matrix of size 1024

Reference execution time: 0 ms

Multi-threaded execution time: 353.456 ms

Performance counter stats for './cmmulti data/input_1024.in':

3,409.69 msec task-clock	# 1.447 CPUs utilized	
15 context-switches	# 0.004 K/sec	
0 cpu-migrations	# 0.000 K/sec	
3,223 page-faults	# 0.945 K/sec	
13,23,97,48,137 cycles	# 3.883 GHz	(31.46%)
31,39,48,17,133 instructions	# 2.37 insn per cycle	(39.23%)
3,19,97,26,565 branches	# 938.422 M/sec	(39.16%)
25,32,269 branch-misses	# 0.08% of all branches	(38.96%)
15,90,49,26,072 L1-dcache-loads	# 4664.626 M/sec	(38.61%)
61,85,87,509 L1-dcache-load-misses	# 3.89% of all L1-dcache accesses	(38.15%)
58,18,50,079 LLC-loads	# 170.646 M/sec	(29.96%)
1,53,815 LLC-load-misses	# 0.03% of all LL-cache accesses	(30.14%)
<not supported>	L1-icache-loads	
10,54,519 L1-icache-load-misses		(31.54%)
15,46,11,14,058 dTLB-loads	# 4534.464 M/sec	(31.75%)
36,496 dTLB-load-misses	# 0.00% of all dTLB cache accesses	(31.81%)
6,854 iTLB-loads	# 0.002 M/sec	(31.64%)

6,721 iTLB-load-misses # 98.06% of all iTLB cache accesses (31.49%)
<not supported> L1-dcache-prefetches
<not supported> L1-dcache-prefetch-misses

2.356414789 seconds time elapsed

3.390231000 seconds user

0.019989000 seconds sys

N=2048,

Input matrix of size 2048

Reference execution time: 0 ms

Multi-threaded execution time: 3697.09 ms

Performance counter stats for './cmmmulti data/input_2048.in':

78,760.88 msec task-clock	#	1.158 CPUs utilized	
4,225 context-switches	#	0.054 K/sec	
41 cpu-migrations	#	0.001 K/sec	
12,440 page-faults	#	0.158 K/sec	
2,99,57,20,28,964 cycles	#	3.804 GHz	(30.81%)
2,47,67,49,79,479 instructions	#	0.83 insn per cycle	(38.50%)
24,69,31,44,278 branches	#	313.520 M/sec	(38.51%)
1,14,23,995 branch-misses	#	0.05% of all branches	(38.49%)
1,26,37,97,49,111 L1-dcache-loads	#	1604.601 M/sec	(38.50%)
7,41,69,35,545 L1-dcache-load-misses	#	5.87% of all L1-dcache accesses	(38.50%)
4,34,57,46,691 LLC-loads	#	55.176 M/sec	(30.78%)
2,77,84,61,799 LLC-load-misses	#	63.94% of all LL-cache accesses	(30.76%)
<not supported> L1-icache-loads			
1,63,02,304 L1-icache-load-misses			(30.76%)
1,25,52,57,24,215 dTLB-loads	#	1593.757 M/sec	(30.75%)
4,30,87,45,277 dTLB-load-misses	#	3.43% of all dTLB cache accesses	(30.75%)
1,37,471 iTLB-loads	#	0.002 M/sec	(30.77%)
2,83,741 iTLB-load-misses	#	206.40% of all iTLB cache accesses	(30.80%)
<not supported> L1-dcache-prefetches			
<not supported> L1-dcache-prefetch-misses			

67.988130952 seconds time elapsed

78.680092000 seconds user

0.084004000 seconds sys

N=4096,

Input matrix of size 4096

Reference execution time: 0 ms

Multi-threaded execution time: 30505.6 ms

Performance counter stats for './cmmsingle data/input_4096.in':

1,20,454.29 msec task-clock	#	3.740 CPUs utilized	
2,021 context-switches	#	0.017 K/sec	
131 cpu-migrations	#	0.001 K/sec	
41,109 page-faults	#	0.341 K/sec	

3,85,44,97,32,089	cycles	#	3.200 GHz	(30.79%)
11,34,81,01,68,136	instructions	#	2.94 insn per cycle	(38.49%)
1,24,11,79,91,978	branches	#	1030.416 M/sec	(38.48%)
2,75,08,455	branch-misses	#	0.02% of all branches	(38.47%)
5,56,22,24,20,375	L1-dcache-loads	#	4617.705 M/sec	(38.46%)
6,75,98,86,320	L1-dcache-load-misses	#	1.22% of all L1-dcache accesses	(38.45%)
8,57,63,385	LLC-loads	#	0.712 M/sec	(30.75%)
6,40,38,576	LLC-load-misses	#	74.67% of all LL-cache accesses	(30.77%)
<not supported>	L1-icache-loads			
3,52,10,479	L1-icache-load-misses			(30.77%)
5,54,78,96,79,265	dTLB-loads	#	4605.811 M/sec	(30.79%)
6,85,03,870	dTLB-load-misses	#	0.01% of all dTLB cache accesses	(30.79%)
7,76,594	iTLB-loads	#	0.006 M/sec	(30.79%)
4,74,956	iTLB-load-misses	#	61.16% of all iTLB cache accesses	(30.80%)
<not supported>	L1-dcache-prefetches			
<not supported>	L1-dcache-prefetch-misses			

32.209103747 seconds time elapsed

120.316759000 seconds user

0.139972000 seconds sys

Conclusion:

So the speedup obtained in single threaded is: ~5.5

And speedup obtained in multi threaded of 4 threads: ~17.28 (From single to multi speed up is 3.14)